

Cultural Data Infrastructure Empirical Audit

This repository contains the supporting materials for:

Haidemariam, T. “Before the Training Data: Cultural Data Infrastructure and the Provenance Structure of Generative AI.”

The accompanying Excel workbook documents the comparative audit of seven publicly documented language-model corpora: C4, The Pile, RedPajama-1T, Dolma, FineWeb, ROOTS, and OSCAR 22.01.

The workbook contains four sheets:

- **Corpus Overview** — corpus characteristics, construction strategies, language coverage, coding, and primary sources.
- **Coding Matrix** — definitions and operational rules for the six cultural-data-infrastructure categories (CI-1–CI-6).
- **Quantitative Shares** — documented component-level token counts and calculated shares for RedPajama-1T and Dolma.
- **Sources & Notes** — primary documentary sources and methodological notes.

Coding records only source categories explicitly identifiable in corpus documentation. A value of 0 therefore means that a category was not separately documented, not that the corresponding type of material is absent from the corpus. NA denotes cases where a comparable classification was not made.

The audit examines documented provenance structure rather than the complete contents of each corpus and does not estimate causal effects of corpus composition on model performance.

All underlying corpus documentation is publicly available through the source links provided in the workbook.